

An introduction to Sequential Monte Carlo

Thang Bui Jes Frelsen

Department of Engineering
University of Cambridge

Research and Communication Club
6 February 2014



UNIVERSITY OF
CAMBRIDGE

Sequential Monte Carlo (SMC) methods

- ▶ Initially designed for online inference in dynamical systems
 - ▶ Observations arrive sequentially and one needs to update the posterior distribution of hidden variables
 - ▶ Analytically tractable solutions are available for linear Gaussian models, but not for complex models
 - ▶ Examples: target tracking, time series analysis, computer vision
- ▶ Increasingly used to perform inference for a wide range of applications, not just dynamical systems
 - ▶ Example: graphical models, population genetic, ...
- ▶ SMC methods are scalable, easy to implement and flexible!

Outline

Motivation

References

Introduction

- MCMC and importance sampling

- Sequential importance sampling and resampling

- Example: A dynamical system

- Proposal

- Smoothing

- MAP estimation

- Parameter estimation

A generic SMC algorithm

Particle MCMC

Particle learning for GP regression

Summary

Bibliography I

Andrieu, Christophe, Doucet, Arnaud, & Holenstein, Roman. 2010.

Particle markov chain monte carlo methods.

Journal of the Royal Statistical Society: Series B (Statistical Methodology), **72**(3), 269–342.

Cappé, Olivier, Godsill, Simon J, & Moulines, Eric. 2007.

An overview of existing methods and recent advances in sequential Monte Carlo.

Proceedings of the IEEE, **95**(5), 899–924.

Doucet, Arnaud, De Freitas, Nando, & Gordon, Neil. 2001.

An introduction to sequential Monte Carlo methods.

Pages 3–14 of: Sequential Monte Carlo methods in practice.
Springer.

Gramacy, Robert B., & Polson, Nicholas G. 2011.

Particle Learning of Gaussian Process Models for Sequential Design and Optimization.

Journal of Computational and Graphical Statistics, **20**(1), 102–118.

Holenstein, Roman. 2009.

Particle Markov Chain Monte Carlo.

Ph.D. thesis, The University Of British Columbia.

Bibliography II

Holenstein, Roman, & Doucet, Arnaud. 2007.

Particle Markov Chain Monte Carlo.

Workshop: New directions in Monte Carlo Methods Fleurance, France.

Wilkinson, Richard. 2014.

GPs huh? what are they good for?

Gaussian Process Winter School, Sheffield.

State Space Models

(Doucet *et al.* , 2001; Cappé *et al.* , 2007)

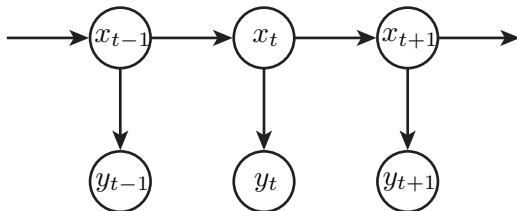
The Markovian, nonlinear, non-Gaussian state space model

- ▶ Unobserved signal or states $\{x_t | t \in \mathbb{N}\}$
- ▶ Observations or output $\{y_t | t \in \mathbb{N}^+\}$ or $\{y_t | t \in \mathbb{N}\}$

$$P(x_0)$$

$$P(x_t | x_{t-1}) \quad \text{for } t \geq 1 \quad (\text{transition probability})$$

$$P(y_t | x_t) \quad \text{for } t \geq 0 \quad (\text{emission/observation probability})$$



Inference for State Space Model

(Doucet *et al.* , 2001; Cappé *et al.* , 2007)

- ▶ We are interested the **posterior distributions** of the unobserved signal

$P(x_{0:t}|y_{0:t})$ – fixed interval smoothing distribution

$P(x_{t-L}|y_{0:t})$ – fixed lag smoothing distribution

$P(x_t|y_{0:t})$ – filtering distribution

- ▶ and **expectations** under these posteriors, e.g.

$$\mathbb{E}_{P(x_{0:t}|y_{0:t})}(h_t) = \int h_t(x_{0:t})P(x_{0:t}|y_{0:t}) \mathrm{d}x_{0:t}$$

for some function $h_t : \mathcal{X}^{(t+1)} \rightarrow \mathbb{R}^{n_{h_t}}$

Couldn't we use *MCMC*?

(Doucet *et al.* , 2001; Holenstein, 2009)

- ▶ Sure, generate N samples from $P(x_{0:t}|y_{0:t})$ using MH
 - ▶ Sample a candidate $x'_{0:t}$ from a proposal distribution

$$x'_{0:t} \sim q(x'_{0:t}|x_{0:t})$$

- ▶ Accept the candidate $x'_{0:t}$ with probability

$$\alpha(x'_{0:t}|x_{0:t}) = \min \left[1, \frac{P(x'_{0:t}|y_{0:t})q(x_{0:t}|x'_{0:t})}{P(x_{0:t}|y_{0:t})q(x'_{0:t}|x_{0:t})} \right]$$

- ▶ Obtain a set of sample $\{x_{0:t}^{(i)}\}_{i=1}^N$
- ▶ Calculate empirical estimates for posterior and expectation

$$\tilde{P}(x_{0:t}|y_{0:t}) = \frac{1}{N} \sum_i \delta_{x_{0:t}^{(i)}}(x_{0:t})$$

$$\mathbb{E}_{\tilde{P}(x_{0:t}|y_{0:t})}(h_t) = \int h_t(x_{0:t}) \tilde{P}(x_{0:t}) \mathrm{d}x_{0:t} = \frac{1}{N} \sum_{i=1}^N h_t(x_{0:t}^{(i)})$$

Couldn't we use *MCMC*?

(Doucet *et al.* , 2001; Holenstein, 2009)

- ▶ Unbiased estimates and in most cases nice convergence

$$\mathbb{E}_{\tilde{P}(x_{0:t}|y_{0:t})}(h_t) \xrightarrow{\text{a.s.}} \mathbb{E}_{P(x_{0:t}|y_{0:t})}(h_t) \quad \text{as } N \rightarrow \infty$$

- ▶ Problem solved!?
- ▶ I can be hard to design a good proposal q
 - ▶ Single-site updates $q(x'_j|x_{0:t})$ can lead to slow mixing
- ▶ What happens if we get a new data point y_{t+1} ?
 - ▶ We cannot (directly) reuse the samples $\{x_{0:t}^{(i)}\}$
 - ▶ We have to run a new MCMC simulations for $P(x_{0:t+1}|y_{0:t+1})$
- ▶ MCMC not well-suited for recursive estimation problems

What about *importance sampling*?

(Doucet *et al.* , 2001)

- ▶ Generate N i.i.d. samples $\{x_{0:t}^{(i)}\}_{i=1}^N$ from an arbitrary **importance sampling distribution** $\pi(x_{0:t}|y_{0:t})$
- ▶ The empirical estimates are

$$\hat{P}(x_{0:t}|y_{0:t}) = \frac{1}{N} \sum_{i=1}^N \delta_{x_{0:t}^{(i)}}(x_{0:t}) \tilde{w}_t^{(i)}$$

$$\mathbb{E}_{\hat{P}(x_{0:t}|y_{0:t})}(h_t) = \frac{1}{N} \sum_{i=1}^N h_t(x_{0:t}^{(i)}) \tilde{w}_t^{(i)}$$

where the **importance weights** are

$$w(x_{0:t}) = \frac{P(x_{0:t}|y_{0:t})}{\pi(x_{0:t}|y_{0:t})} \quad \text{and} \quad \tilde{w}_t^{(i)} = \frac{w(x_{0:t}^{(i)})}{\sum_j w(x_{0:t}^{(j)})}$$

What about *importance sampling*?

(Doucet *et al.* , 2001)

- ▶ $\mathbb{E}_{\hat{P}(x_{0:t}|y_{0:t})}(h_t)$ is biased, but converges to $\mathbb{E}_{P(x_{0:t}|y_{0:t})}(h_t)$
- ▶ **Problem solved!?**
- ▶ Designing a good importance distribution can be hard!
- ▶ Still not adequate for **recursive estimation**
 - ▶ When seeing new data y_{t+1} , we **cannot reuse** the samples and weights for time t

$$\{x_{0:t}^{(i)}, \tilde{w}_t^{(i)}\}_{i=1}^N$$

to sample from $P(x_{0:t+1}|y_{0:t+1})$

Sequential importance sampling

(Doucet *et al.* , 2001; Cappé *et al.* , 2007)

- Assume that the **importance distribution** can be factored as

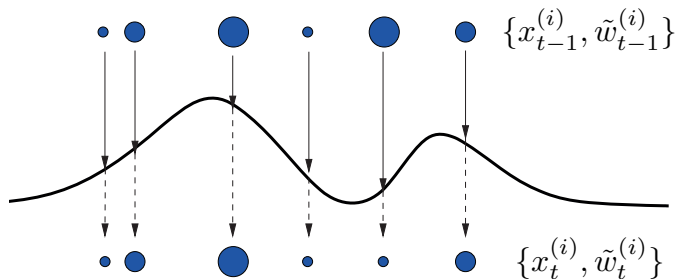
$$\begin{aligned}\pi(x_{0:t}|y_{0:t}) &= \underbrace{\pi(x_{0:t-1}|y_{0:t-1})}_{\text{importance distribution at time } t-1} \underbrace{\pi(x_t|x_{0:t-1}, y_{0:t})}_{\text{extension to time } t} \\ &= \pi(x_0|y_0) \prod_{k=1}^t \pi(x_k|x_{0:k-1}, y_{0:k})\end{aligned}$$

- The importance weight can then be evaluated recursively

$$\tilde{w}_t^{(i)} \propto \tilde{w}_{t-1}^{(i)} \frac{P(y_t|x_t^{(i)})P(x_t^{(i)}|x_{t-1}^{(i)})}{\pi(x_t^{(i)}|x_{0:t-1}, y_{0:t})} \quad (1)$$

- Given past i.i.d. trajectories $\{x_{0:t-1}^{(i)} | i = 1, \dots, N\}$ we can
 - simulate $x_t^{(i)} \sim \pi(x_t|x_{0:t-1}^{(i)}, y_{0:t})$
 - update the weight $\tilde{w}_t^{(i)}$ for $x_{0:t}^{(i)}$ based on $\tilde{w}_{t-1}^{(i)}$ using eq. (1)
- Note that the extended trajectories $\{x_{0:t}^{(i)}\}$ remain i.i.d.

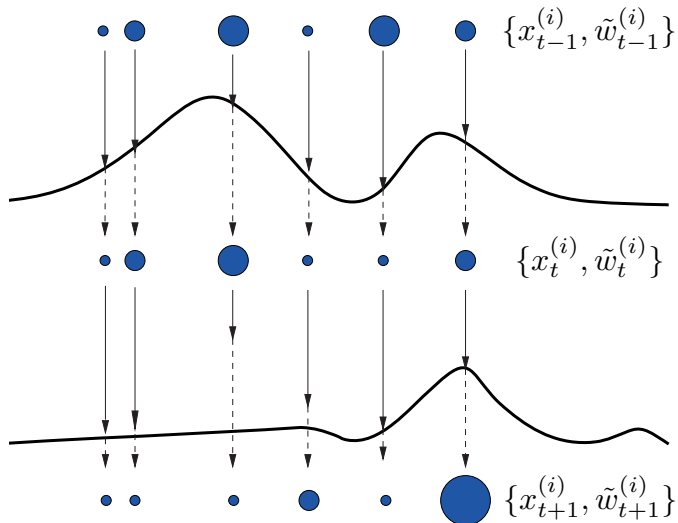
Sequential importance sampling



Adapted from (Doucet *et al.* , 2001)

► Problem solved!?

Sequential importance sampling



Adapted from (Doucet *et al.*, 2001)

- Weights become **highly degenerated** after few steps

Sequential importance resampling

(Doucet *et al.* , 2001; Cappé *et al.* , 2007)

- ▶ Key idea to eliminate weight degeneracy
 1. **Eliminate** particles with low importance weights
 2. **Multiply** particles with high importance weights
- ▶ Introduce a resampling each time step (or “occasionally”)
- ▶ Resample a new trajectory $\{x_{0:t}'^{(i)} | i = 1, \dots, N\}$
 - ▶ Draw N samples from

$$\hat{P}(x_{0:t}|y_{0:t}) = \frac{1}{N} \sum_{i=1}^N \delta_{x_{0:t}^{(i)}}(x_{0:t}) \tilde{w}_t^{(i)}$$

- ▶ The weights of the new samples are $\tilde{w}_t'^{(i)} = \frac{1}{N}$
- ▶ The new empirical (unweighted) distribution a time step t

$$\hat{P}'(x_{0:t}|y_{0:t}) = \frac{1}{N} \sum_{i=1}^N \delta_{x_{0:t}^{(i)}}(x_{0:t}) N_t^{(i)}$$

where $N_t^{(i)}$ is the number of copies of $x_{0:t}^{(i)}$.

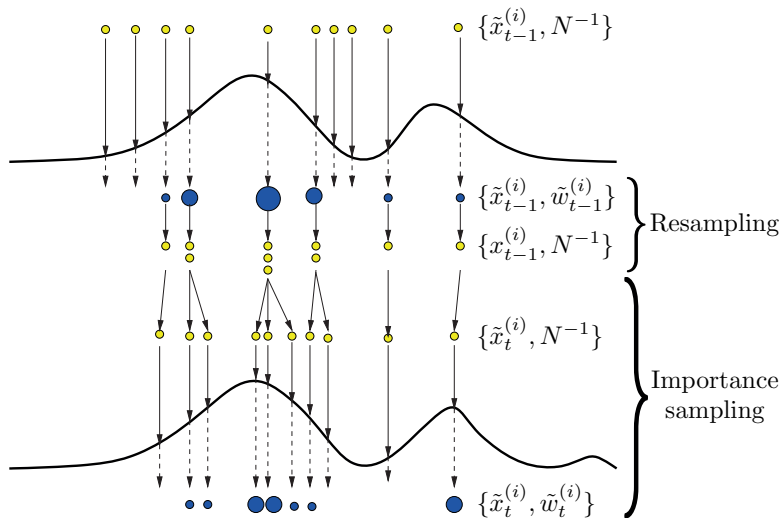
- ▶ $N_t^{(i)}$ is sampled for a multinomial with parameters $w_t^{(i)}$

Sequential importance resampling

(Doucet *et al.* , 2001; Cappé *et al.* , 2007)

- 1: **for** $i = 1, \dots, N$ **do**
- 2: Sample $x_0^{(i)} \sim \pi(x_0|y_0)$
- 3: $w_0^{(i)} \leftarrow \frac{P(y_0|x_0^{(i)})P(x_0^{(i)})}{\pi(x_0^{(i)}|y_0)}$
- 4: **for** $t = 1, \dots, T$ **do**
 Importance sampling step
- 5: **for** $i = 1, \dots, N$ **do**
- 6: Sample $\tilde{x}_t^{(i)} \sim \pi(x_t|x_{0:t-1}, y_{0:t})$
- 7: $\tilde{x}_{0:t}^{(i)} \leftarrow (x_{0:t-1}^{(i)}, \tilde{x}_t^{(i)})$
- 8: $\tilde{w}_t^{(i)} \leftarrow w_{t-1}^{(i)} \frac{P(y_t|x_t^{(i)})P(x_t^{(i)}|x_{t-1}^{(i)})}{\pi(x_t^{(i)}|x_{0:t-1}^{(i)}, y_{0:t})}$
- Resampling/selection step
- 9: Sample N particles $\{x_{0:t}^{(i)}\}$ from $\{\tilde{x}_{0:t}^{(i)}\}$ according to $\{\tilde{w}_t^{(i)}\}$
- 10: $w_t^{(i)} \leftarrow \frac{1}{N}$ **for** $i = 1, \dots, N$
- 11: **return** $\{x_{0:t}^{(i)}\}_{i=1}^N$

Sequential importance resampling



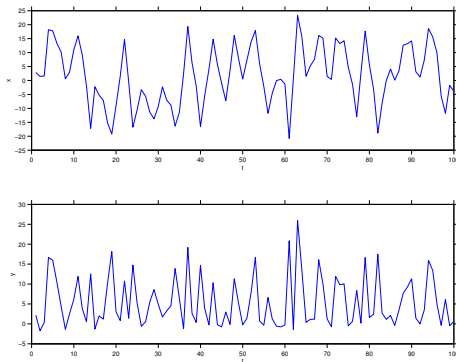
$i = 1, \dots, N$ and $N = 10$, figure modified from (Doucet *et al.*, 2001)

Example - A dynamical system

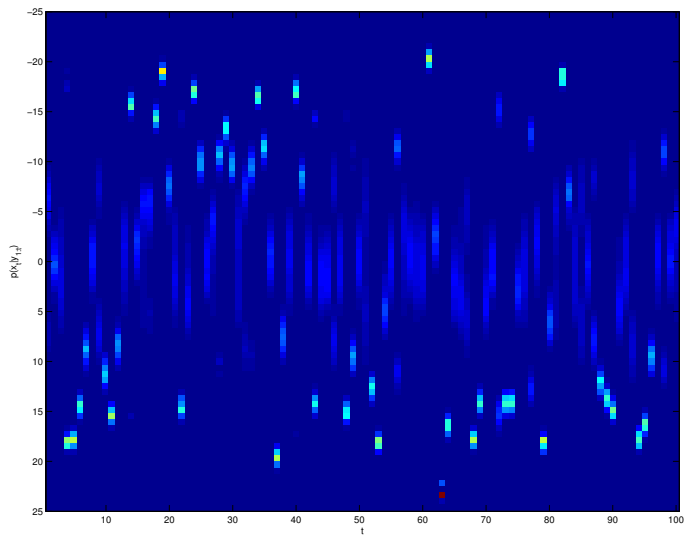
$$\begin{aligned}x_t &= \frac{1}{2}x_{t-1} + 25\frac{x_{t-1}}{1+x_{t-1}^2} + 8\cos(1.2t) + u_t \\y_t &= \frac{x_t^2}{20} + v_t\end{aligned}$$

where

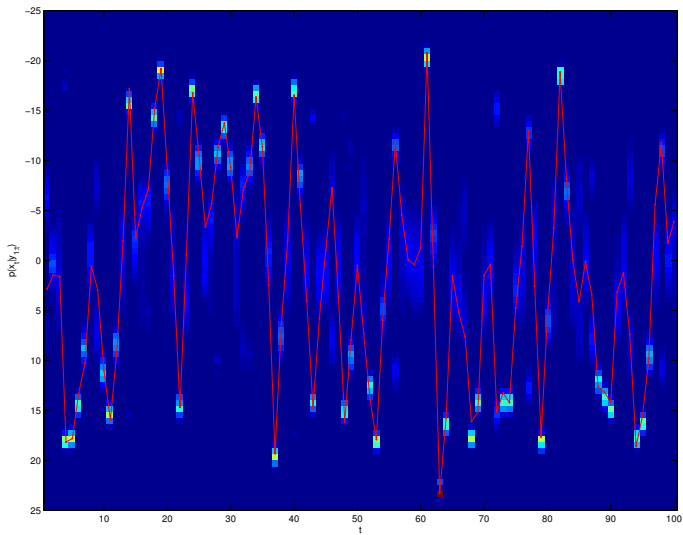
$x_0 \sim \mathcal{N}(0, \sigma_0^2)$, $u_t \sim \mathcal{N}(0, \sigma_u^2)$, $v_t \sim \mathcal{N}(0, \sigma_v^2)$, $\sigma_0^2 = \sigma_u^2 = 10$, $\sigma_v^2 = 1$.



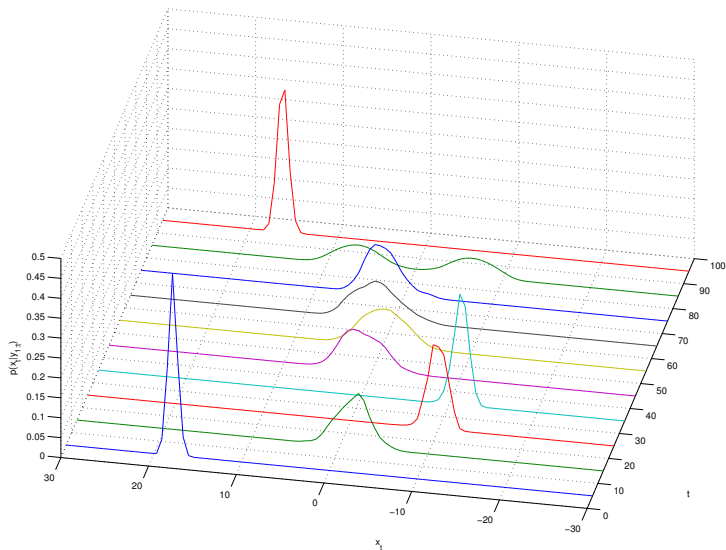
Posterior distribution of states



Posterior distribution of states



Posterior distribution of states



Proposal

- ▶ **Bootstrap filter** uses $\pi_t(x_t|x_{t-1}, y_t) = P(x_t|x_{t-1})$ which leads to a simple form for the importance weight update:
 $w_t^{(i)} \propto w_{t-1}^{(i)} P(y_t|x_t^{(i)})$
 - ▶ The weight update depends on the new **proposed state** and the observation!
 - ▶ Uninformative observation can lead to poor performance
- ▶ Optimal proposal:

$$\pi_t(x_t|x_{t-1}, y_t) = P(x_t|x_{t-1}, y_t)$$

therefore: $w_t^{(i)} \propto w_{t-1}^{(i)} P(y_t|x_{t-1}) = \int P(y_t|x_t) P(x_t|x_{t-1}) dx_t$

- ▶ The weight update depends on the **previous state** and the observation
- ▶ Analytically intractable integral, need to resort to approximation techniques.

Smoothing

- ▶ For a complex SSM, the posterior distribution of state variables can be *smoothed* by including future observations.
- ▶ The joint smoothing distribution can be factorised:

$$\begin{aligned} P(x_{0:T}|y_{0:T}) &= P(x_T|y_{0:T}) \prod_{t=0}^{T-1} P(x_t|x_{t+1}, y_{0:t}) \\ &\propto P(x_T|y_{0:T}) \prod_{t=0}^{T-1} \underbrace{P(x_t|y_{0:t})}_{\text{filtering distribution}} \underbrace{P(x_{t+1}|x_t)}_{\text{likelihood of future state}} \end{aligned}$$

Hence, the weight update:

$$\hat{w}_t^{(i)}(x_{t+1}) = w_t^{(i)} P(x_{t+1}|x_t)$$

Particle smoother

Algorithm:

- ▶ Run forward simulation to obtain particle paths $\{x_t^{(i)}, w_t^{(i)}\}_{i=1, \dots, N; t=1, \dots, T}$
- ▶ Draw \tilde{x}_T from $\hat{P}(x_T|y_{0:T})$
- ▶ Repeat:
 - ▶ Adjust and normalise the filtering weights $w_t^{(i)}$:

$$\hat{w}_t^{(i)} = w_t^{(i)} P(\tilde{x}_{t+1}|x_t)$$

- ▶ Draw a random sample \tilde{x}_t from $\hat{P}(x_{t:T}|y_{0:T})$

The sequence $(\tilde{x}_0, \tilde{x}_2, \dots, \tilde{x}_T)$ is a random draw from the approximate distribution $\hat{P}(x_{0:T}|y_{0:T})$ [$\mathcal{O}(NT)$]

MAP estimation

- ▶ Maximum a posteriori (MAP) estimate:

$$\operatorname{argmax}_{x_{0:T}} P(x_{0:T}|y_{0:T}) = \operatorname{argmax}_{x_{0:T}} P(x_0) \prod_{t=1}^T P(x_t|x_{t-1}) \prod_{t=0}^T P(y_t|x_t)$$

Question: Can we just choose particle trajectory with largest weights?

Answer: **NO!**

- ▶ Assume a discrete particle grid, $x_t \in x_t^{(i)}_{1 \leq i \leq N}$, the approximation can be interpreted as a **Hidden Markov Model** with N states.

MAP estimate can be found using the Viterbi algorithm

- ▶ Keep track of the probability of the most likely path so far
- ▶ Keep track of the last state index of the most likely path so far

Viterbi algorithm for MAP estimation

Observation

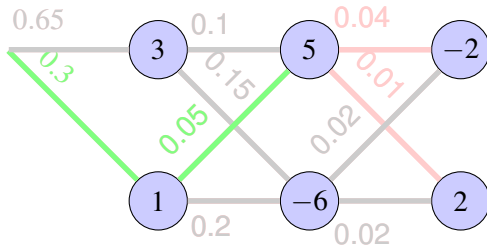
-1

4

2

Particle 1

Particle 2



Path probability update:

$$\alpha_t^{(j)} = \alpha_{t-1}^{(j)} P(x_t^{(i)} | x_{t-1}^{(i)}) P(y_t | x_t^{(i)})$$

Parameter estimation

- ▶ Consider SSMs that have $P(x_t|x_{t-1}, \theta)$, $P(y_t|x_t, \theta)$ where θ is a static parameter vector and one wishes to estimate θ .
- ▶ Marginal likelihood:

$$l(y_{0:T}|\theta) = \int p(y_{0:T}, x_{0:T}|\theta) dx_{0:T}$$

- ▶ Optimise $l(y_{0:T}|\theta)$ using the EM algorithm:
 - ▶ E-step:

$$\hat{\tau}(\theta, \theta_k) = \sum_{i=1}^N w_T^{(i, \theta)} \sum_{t=0}^{T-1} s_{t, \theta}(x_t^{(i, \theta_k)}, x_{t+1}^{(i, \theta_k)})$$

where

$$s_{t, \theta}(x_t, x_{t+1}) = \log(P(x_{t+1}|x_t, \theta)) + \log(P(y_{t+1}|x_{t+1}, \theta))$$

- ▶ Optimise $\hat{\tau}(\theta|\theta_k)$ to update θ_k .

A generic SMC algorithm

(Holenstein, 2009)

- ▶ We want to sample from a **target distribution** $\pi(\mathbf{x})$, $\mathbf{x} \in \mathcal{X}^p$
- ▶ Assume we have a sequence of **bridging distributions** of **increasing dimension**

$$\{\pi_n(\mathbf{x}_n)\}_{n=1}^p = \{\pi_1(x_1), \pi_2(x_1, x_2), \dots, \pi_p(x_1, \dots, x_p)\}$$

where

- ▶ $\pi_n(\mathbf{x}_n) = Z_n^{-1} \gamma_n(\mathbf{x}_n)$
- ▶ $\pi_p(\mathbf{x}_p) = \pi(\mathbf{x})$
- ▶ A sequence of **importance densities** on \mathcal{X}

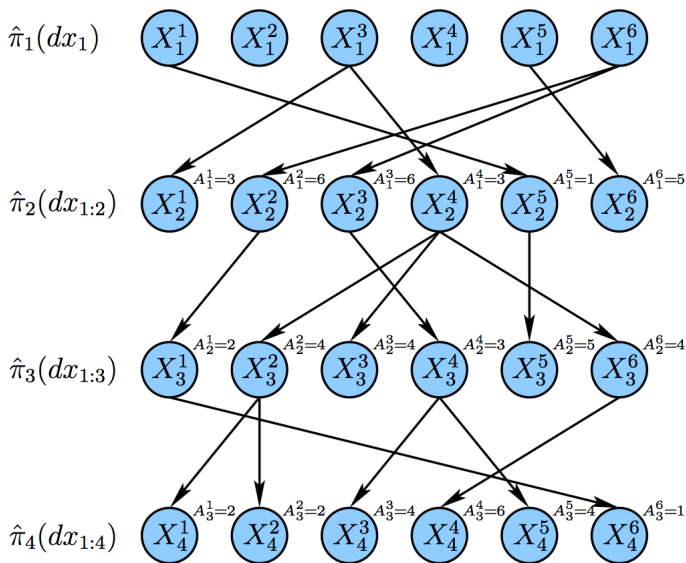
$$\underbrace{M_1(x_1)}_{\text{for initial sample}}, \quad \underbrace{\{M_n(x_n | \mathbf{x}_{n-1})\}_{n=2}^p}_{\substack{\text{for extending } \mathbf{x}_{n-1} \in \mathcal{X}^{n-1} \\ \text{by sampling } x_n \in \mathcal{X}}}$$

- ▶ A **resampling distribution**

$$r(A_n | \mathbf{w}_n), \quad A_n \in \{1, \dots, N\}^N \text{ and } \mathbf{w}_n \in [0, 1]^N$$

where A_{n-1}^i is the parent at time $n-1$ of some particle \mathbf{X}_n^i

A generic SMC algorithm



From (Holenstein, 2009)

A generic SMC algorithm

1 **At** $n = 1$

2 Sample $\mathbf{X}_1^i \sim M_1(\cdot)$

3 Update and normalise the weights

$$w_1(\mathbf{X}_1^i) := \frac{\gamma_1(\mathbf{X}_1^i)}{M_1(\mathbf{X}_1^i)}, \quad W_1^i = \frac{w_1(\mathbf{X}_1^i)}{\sum_{k=1}^N w_1(\mathbf{X}_1^k)}.$$

4 **For** $n = 2, \dots, p$ **do**

5 Sample $\mathbf{A}_{n-1} \sim r(\cdot | \mathbf{W}_{n-1})$

6 Sample $X_n^i \sim M_n(\mathbf{X}_{n-1}^{\mathbf{A}_{n-1}^i}, \cdot)$ and set $\mathbf{X}_n^i = (\mathbf{X}_{n-1}^{\mathbf{A}_{n-1}^i}, X_n^i)$

7 Update and normalise the weights

$$w_n(\mathbf{X}_n^i) := \frac{\gamma_n(\mathbf{X}_n^i)}{\gamma_{n-1}(\mathbf{X}_{n-1}^{\mathbf{A}_{n-1}^i}) M_n(\mathbf{X}_{n-1}^{\mathbf{A}_{n-1}^i}, X_n^i)},$$

$$W_n^i = \frac{w_n(\mathbf{X}_n^i)}{\sum_{k=1}^N w_n(\mathbf{X}_n^k)}$$

A generic SMC algorithm

(Holenstein, 2009)

- ▶ Again we can calculate empirical estimates for **target** and the **normalization constant** ($\pi(\mathbf{x}) = Z^{-1}\gamma(\mathbf{x})$)

$$\hat{\pi}^N(\mathbf{x}) = \sum_{i=1}^N \delta_{\mathbf{x}_p^{(i)}}(\mathbf{x}) W_p^{(i)}$$
$$\hat{Z}^N(\mathbf{x}) = \prod_{n=1}^p \left(\frac{1}{N} \sum_{i=1}^N w_n(X_n^{(i)}) \right)$$

- ▶ Convergence can be shown under weak assumptions

$$\hat{\pi}^N(\mathbf{x}) \xrightarrow{\text{a.s.}} \pi(\mathbf{x}) \quad \text{as } N \rightarrow \infty$$
$$\hat{Z}^N \xrightarrow{\text{a.s.}} Z \quad \text{as } N \rightarrow \infty$$

- ▶ The SIR algorithm for **state space models** is a **special case** of this generic SMC algorithm

Motivation for Particle MCMC

(Holenstein & Doucet, 2007; Holenstein, 2009)

- ▶ Let's return to the problem of sampling from a target

$$\pi(\mathbf{x}), \text{ where } \mathbf{x} = (x_1, x_2, \dots, x_n) \in \mathcal{X}^n$$

using MCMC

- ▶ Single-site proposal $q(x'_j|\mathbf{x})$
 - ▶ Easy to design
 - ▶ Often leads to **slow mixing**
- ▶ It would be more efficient, if we could update larger blocks
 - ▶ Such proposals are harder to construct
- ▶ We could use SMC as a proposal distribution!

Particle Metropolis Hastings Sampler

1 **Initialisation** $i = 0$

2 Run an SMC algorithm targeting $\pi(\mathbf{x})$

3 sample $\mathbf{X}(0) \sim \hat{\pi}^N(\cdot)$ and compute $\hat{Z}^N(0)$

4 **For iteration** $i \geq 1$

5 Run an SMC algorithm targeting $\pi(\mathbf{x})$, sample $\mathbf{X}^* \sim \hat{\pi}^N(\cdot)$ and
compute $\hat{Z}^{N,*}$

6 With probability

$$1 \wedge \frac{\hat{Z}^{N,*}}{\hat{Z}^N(i-1)}, \quad (3.3)$$

set $\mathbf{X}(i) = \mathbf{X}^*$ and $\hat{Z}^N(i) = \hat{Z}^{N,*}$, otherwise set $\mathbf{X}(i) = \mathbf{X}(i-1)$ and
 $\hat{Z}^N(i) = \hat{Z}^N(i-1)$

From (Holenstein, 2009)

Particle Metropolis Hastings (PMH) Sampler

(Holenstein, 2009)

- ▶ Standard **independent** MH algorithm ($q(x'|x) = q(x')$)
- ▶ Target $\tilde{\pi}^N$ and proposal q^N defined on an extended space

$$\frac{\tilde{\pi}^N(\cdot)}{q^N(\cdot)} = \frac{\hat{Z}^N}{Z}$$

which leads to the acceptance ratio

$$\alpha = \min \left[1, \frac{\hat{Z}^{N,*}}{\hat{Z}^N(i-1)} \right]$$

- ▶ Note that $\alpha \rightarrow 1$ as $N \rightarrow \infty$, since $\hat{Z}^N \rightarrow Z$ as $N \rightarrow \infty$

Particle Gibbs (PG) Sampler

(Holenstein, 2009)

- ▶ Assume that we are interested in sampling from

$$\pi(\theta, \mathbf{x}) = \frac{\gamma(\theta, \mathbf{x})}{Z}$$

- ▶ Assume that sampling from
 - ▶ $\pi(\theta|\mathbf{x})$ is **easy**
 - ▶ $\pi(\mathbf{x}|\theta)$ is **hard**
- ▶ The PG Sampler uses SMC to sample from $\pi(\mathbf{x}|\theta)$
 1. Sample $\theta(i) \sim \pi(\theta|\mathbf{x}(i-1))$
 2. Sample $\mathbf{x}(i) \sim \hat{\pi}^N(\mathbf{x}|\theta(i))$
- ▶ If sampling from $\pi(\theta|\mathbf{x})$ is not easy?
 - ▶ We can use a MH update for θ

Parameter estimation a state space models using PG

(Andrieu *et al.* , 2010)

- ▶ (Re)consider the non-linear state space model

$$x_t = \frac{1}{2}x_{t-1} + 25\frac{x_{t-1}}{1 + x_{t-1}^2} + 8\cos(1.2t) + V_t$$

$$y_t = \frac{x_t^2}{20} + W_t$$

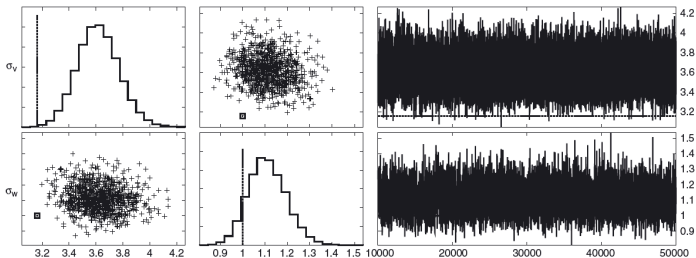
where $x_0 \sim \mathcal{N}(0, \sigma_0^2)$, $V_t \sim \mathcal{N}(0, \sigma_V^2)$ and $W_t \sim \mathcal{N}(0, \sigma_W^2)$

- ▶ Assume that $\theta = (\sigma_V^2, \sigma_W^2)$ is unknown
- ▶ Simulate $y_{1:T}$ for $T = 500$, $\sigma_0^2 = 5$, $\sigma_V^2 = 10$ and $\sigma_W^2 = 1$
- ▶ Sample from $P(\theta, x_{1:t}|y_{1:t})$ using
 - ▶ Particle Gibbs sampler, with importance dist. $f_\theta(x_n|x_{n-1})$
 - ▶ One-at-a-time MH sampler, with proposal $f_\theta(x_n|x_{n-1})$
- ▶ The algorithms ran for 50,000 iterations (burn-in of 10,000)
 - ▶ Vague inverse-Gamma priors for $\theta = (\sigma_V^2, \sigma_W^2)$

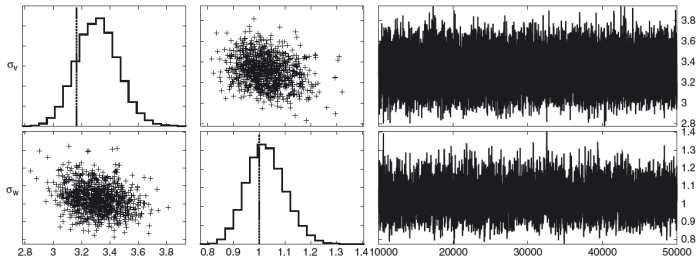
Parameter estimation a state space models using PG

(Andrieu *et al.* , 2010)

One-at-a-time Metropolis Hastings



Particle Gibbs Sampler



Particle learning for GP regression – motivation

Training a GP using data: $\mathcal{D}_{1:n} = \{(x_1, y_1), \dots, (x_n, y_n)\}$ and make prediction:

$$P(y^* | \hat{\theta}, \mathcal{D}, x^*) \quad (2)$$

or

$$P(y^* | \mathcal{D}, x^*) = \int P(y^* | \theta, \mathcal{D}, x^*) P(\theta | \mathcal{D}) d\theta \quad (3)$$

- ▶ Estimate model hyperparameters θ_n using ML (2) or use sampling to find the posterior distribution (3)
- ▶ Find the inverse of the covariance matrix K_n^{-1} .
- ▶ Computational cost $\mathcal{O}(n^3)$.

Sequential update

Given a new observation pair (x_{n+1}, y_{n+1}) that we want to use in our training set, need to find K_{n+1}^{-1} and re-estimate hyperparameters θ_{n+1} .

- ▶ a naive implementation costs $\mathcal{O}(n^3)$
- ▶ need an efficient approach that makes use of the sequential nature of data.

Particle learning for GP regression

(Gramacy & Polson, 2011; Wilkinson, 2014)

Sufficient information for each particle $S_n^{(i)} = \{K_n^{(i)}, \theta_n^{(i)}\}$

Two-step update based on:

$$\begin{aligned} P(S_{n+1} | \mathcal{D}_{1:n+1}) &= \int P(S_{n+1} | S_n, \mathcal{D}_{n+1}) P(S_n | \mathcal{D}_{1:n+1}) dS_n \\ &\propto \int P(S_{n+1} | S_n, \mathcal{D}_{n+1}) P(\mathcal{D}_{n+1} | S_n) P(S_n | \mathcal{D}_{1:n}) dS_n \end{aligned}$$

1. **Resample** indices $\{i\}_{i=1}^N$ with replacement to obtain new indices $\{\zeta(i)\}_{i=1}^N$ according to weights

$$w_i \sim P(\mathcal{D}_{n+1} | S_n^{(i)}) = P(y_{n+1} | x_{n+1}, \mathcal{D}_n, \theta_n^{(i)})$$

2. **Propagate** sufficient information from S_n to S_{n+1}

$$S_{n+1}^{(i)} \sim P(S_{n+1} | S_n^{\zeta(i)}, \mathcal{D}_{1:n+1})$$

Propagation

- ▶ Parameters θ_n are static and can be deterministically copied from $S_n^{\zeta(i)}$ to $S_{n+1}^{(i)}$.
- ▶ Covariance matrix rank-one update to build K_{n+1}^{-1} from K_n^{-1} :

$$K_{n+1} = \begin{bmatrix} K_n & k(x_{n+1}) \\ k^\top(x_{n+1}) & k(x_{n+1}, x_{n+1}) \end{bmatrix}$$

then

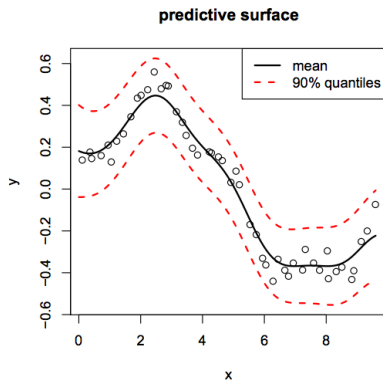
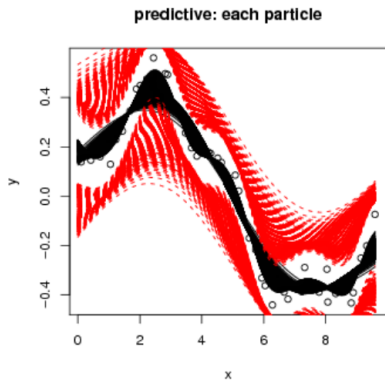
$$K_{n+1}^{-1} = \begin{bmatrix} K_n^{-1} + g_n(x_{n+1})g_n^\top(x_{n+1})/\mu_n(x_{n+1}) & g_n(x_{n+1}) \\ g_n^\top(x_{n+1}) & \mu_n(x_{n+1}) \end{bmatrix}$$

where

$$\begin{aligned} g_n(x) &= -\mu(x)K_n^{-1}k(x) \\ \mu_n(x) &= [k(x, x) - k^\top(x)K_n^{-1}k(x)]^{-1} \end{aligned}$$

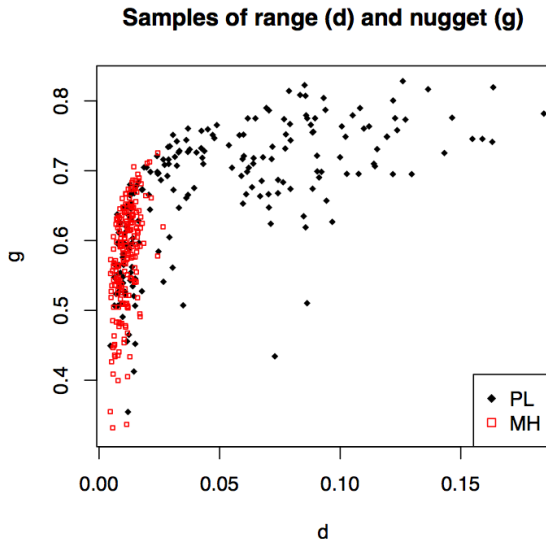
- ▶ Use Cholesky update for stability
- ▶ Cost: $\mathcal{O}(n^2)$

Illustrative result 1 - Prediction



From (Gramacy & Polson, 2011)

Illustrative result 2 - Particle locations



From (Gramacy & Polson, 2011)

SMC for learning GP models

Advantages:

- ▶ Fast for sequential learning problems

Disadvantages:

- ▶ Particle degeneracy/depletion
 - ▶ Use MCMC sampler to augment the propagate and *rejuvenate* the particles after m sequential updates.
- ▶ The predictive distribution given model hyperparameters needs to be *analytically tractable* [See resample step]

Similar treatment for classification can be found in (Gramacy & Polson, 2011).

Summary

- ▶ SMC is a **powerful method** for sampling from distributions with **sequential nature**
 - ▶ **Online learning** in state space models
 - ▶ Sample from **high dimensional** distributions
 - ▶ As **proposal distribution** in MCMC
- ▶ We presented two concrete examples of using SMC
 - ▶ **Particle Gibbs** for sampling from the posterior distributions of the parameters in a non-linear state space model
 - ▶ **Particle learning** of the hyperparameters in a GP model
- ▶ Thank you for your attention!